

Repository logic tree and auditable evidence handoffs

Seven ordered stages connect versioned Agent contracts, explicit execution files and durable intermediate results without allowing presentation artifacts to alter model scoring.

2,365		20		3		50	
literature records		recommended models		test datasets		decision rounds	
PROJECT TREE		AGENT CONTRACTS		EXECUTION		PERSISTED EVIDENCE	
zss/ repository root							
S0	ENTRY & CONFIG	No role-specific Agent Global controls and guards	run_menu.bat -> amp_benchmark_menu.py config.py + .env workflow_guards.py run_manifest.py	runtime selection + validated configuration run checkpoint and artifact provenance			
S1	LITERATURE SEARCH & RECOMMENDATION	agents/shared/ + agents/deepseek_meeting/ evidence_compressor info_extractor model_dataset metric critic -> chief (.md)	agent_md_loader.py deep_research_literature_agent.py dataset_recommendation_agent.py benchmark_portfolio.py	evidence_pool.json -> compact_evidence_pool.json -> screening_decisions.csv -> literature_memory.json/.md dataset_candidate_pool.json -> dataset_plan.json exports/: recommended_models / datasets / metrics.csv	VERIFIED HANDOFF 2,365 papers -> 20 models / 3 datasets / 44 metric rows		
S2	MODEL REGISTRATION & HPC DEPLOYMENT	agents/model_onboarding/ repository_inspector_system.md repository_inspector_task.md agents/model_execution/ self_heal_system.md self_heal_task.md	new_model_onboarding.py -> hpc_model_ops.py README -> registry -> upload -> Conda/deps -> smoke/self-heal	data/local_registry.json repository + environment + command + status data/hpc_self_heal/ failure logs + bounded repair records	VERIFIED HANDOFF registered model state		
S3	DATASET BENCHMARKING	agents/runtime_prompts/ benchmark_architect coder critic pi (.md) dataset_etf_agent.md	AUTO main.py -> run_meeting.py -> workflow_utils.py MANUAL import_manual_prediction_results.py COMMON -> scientific_evaluation.py	data/runs/(run_id)/ manifest + artifacts + results/(dataset)/ data/results_manual/(dataset)/ standardized_predictions.csv + eval_result.json scientific_evaluation.json/.md + Critic report	VERIFIED HANDOFF 3 datasets / 62,310 sequences / 18 usable model outputs		
S4	50-ROUND MULTI-AGENT RANKING	agents/weight_meeting/: shared_system.md literature statistics screening -> reviewer -> chief research_advisor.md	codex_agent_weight_meeting.py or llm_agent_weight_meeting.py -> iterative_weight_meeting.py blind proposals -> audit -> Chief -> 50 bounded updates	agent_evidence_bundle -> three expert proposals -> reviewer_agent_audit -> chief_initial_decision -> rounds/round_001 ... round_050.json -> metric weights / model scores -> final ranking + report	VERIFIED HANDOFF 12 metrics / 50 rounds / 600 weights / 900 scores		
S5	TOP-3 & ENSEMBLE LEARNING	Consumes Stage 4 Chief decision Research Advisor synthesis	ensemble_top3_selector.py amp_research_advisor.py	ensemble_top3_selection.json ensemble_top3_combination_ranking.csv	VERIFIED HANDOFF Top-3 and complementarity		
S6	PAPER FIGURES & QUALITY ASSURANCE	Presentation only; never feeds values back into scoring	figures/ plotting scripts tests/	SVG / PDF / PNG / TIFF + QA notes + source CSV			

Repository invariant

Automatic HPC evaluation and manual prediction import remain separate until the common scientific evaluator; figures and tests never feed values back into scoring.