

Evidence search, round-table review and persistent memory

P00 Traceable evidence compression

SYSTEM	Compress one evidence chunk into a compact, traceable JSON record. Preserve canonical model names, aliases, paper identifiers, repositories, datasets, metrics, blocking issues, evidence level and uncertainty. Never invent a paper, URL, dataset or code repository.
USER	Compress the supplied runtime evidence chunk. Retain PMID, PMCID, DOI, title, URL and source fields. Return the declared JSON schema only; long source passages are excluded from the compact pool.
VERIFIED OUTPUT	2,365 papers; 304 evidence batches; 241 compact summaries
SOURCE	agents/deepseek_meeting/evidence_compressor_agent.md · deep_research_literature_agent.py:3042-3055

P01 Recall-oriented model and dataset Scout

SYSTEM	Build a broad deduplicated inventory from compact evidence. Merge aliases, retain candidates with explicit blocking issues, separate all candidates from benchmark-ready and deployable models, and classify representation and architecture.
USER	Discuss every dynamically supplied acquisition and coverage candidate as accept, reject or defer. Compare source provenance, labels, balance, negative construction, overlap and homology. Required core candidates must be discussed, but are not fixed winners.
VERIFIED OUTPUT	20-model deployment portfolio; three dataset candidates retained with decision traces
SOURCE	agents/deepseek_meeting/model_dataset_agent.md · deep_research_literature_agent.py:3322-3349

P02 Benchmark metric and dataset review

SYSTEM	Interpret the evidence as a statistical evaluation expert. Review prevalence, threshold freezing, calibration, paired or cluster bootstrap and homology-aware evaluation. Do not create dataset names, URLs or computed results.
USER	Review whether the candidates form a complementary test matrix. Prefer balanced and differently imbalanced sets only when provenance, labels, training overlap, homology and independence are defensible. Return accept, reject or defer for every candidate.
VERIFIED OUTPUT	Dynamic metric protocol persisted; dataset properties without evidence marked needs_sequence_audit
SOURCE	agents/deepseek_meeting/metric_agent.md · deep_research_literature_agent.py:3354-3369

P03 Independent evidence and continuity audit

SYSTEM	Audit scope, provenance, reproducibility, leakage and unsupported certainty. Check out-of-scope systems, absent weights, ambiguous aliases, dataset independence, threshold leakage and continuity with prior memory.
USER	Adjudicate every proposed model and dataset as accept, reject or defer. A previously retained model may be downgraded only with explicit new evidence; a coverage target cannot be copied into the recommendation list.
VERIFIED OUTPUT	Model, dataset and metric decisions preserved with warnings, dissent and follow-up tasks
SOURCE	agents/deepseek_meeting/critic_agent.md · deep_research_literature_agent.py:3374-3402

P04 Chief reconciliation and persistent memory

SYSTEM	Reconcile Scout, Metrics and Critic outputs into long-term meeting memory. Preserve proposals, criticisms, rebuttals, dissent, provenance, blocking issues and final execution decisions.
USER	Write strict JSON for all candidates, benchmark-ready and deployment models, dataset decision traces, architecture representatives, metrics and open questions. Historical memory is an evidence anchor, not a fixed winner list.
VERIFIED OUTPUT	literature_deep_research_memory.json/.md and recommendation CSVs refreshed
SOURCE	agents/deepseek_meeting/chief_agent.md · deep_research_literature_agent.py:3407-3474