

Evidence search, memory and literature consensus

P00 Traceable evidence compression

SYSTEM	Compress one evidence chunk into a compact, traceable JSON record. Preserve canonical model names, aliases, paper identifiers, repositories, datasets, metrics, blocking issues, evidence level and uncertainty. Never invent a paper, URL, dataset or code repository.
USER	Compress the supplied runtime evidence chunk. Retain PMID, PMCID, DOI, title, URL and source fields. Return the declared JSON schema only; long source passages are excluded from the compact pool.
VERIFIED OUTPUT	2,365 papers; 304 evidence batches; 241 compact summaries
SOURCE	agents/deepseek_meeting/evidence_compressor_agent.md · deep_research_literature_agent.py:3042-3055

P01 Recall-oriented model and dataset Scout

SYSTEM	Build a broad deduplicated inventory from compact evidence. Merge aliases, retain candidates with explicit blocking issues, separate all candidates from benchmark-ready and deployable models, and classify representation and architecture.
USER	Discuss every dynamically supplied acquisition and coverage candidate as accept, reject or defer. Compare source provenance, labels, balance, negative construction, overlap and homology. Required core candidates must be discussed, but are not fixed winners.
VERIFIED OUTPUT	20-model deployment portfolio; three dataset candidates retained with decision traces
SOURCE	agents/deepseek_meeting/model_dataset_agent.md · deep_research_literature_agent.py:3322-3349

P02 Benchmark metric and dataset review

SYSTEM	Interpret the evidence as a statistical evaluation expert. Review prevalence, threshold freezing, calibration, paired or cluster bootstrap and homology-aware evaluation. Do not create dataset names, URLs or computed results.
USER	Review whether the candidates form a complementary test matrix. Prefer balanced and differently imbalanced sets only when provenance, labels, training overlap, homology and independence are defensible. Return accept, reject or defer for every candidate.
VERIFIED OUTPUT	Dynamic metric protocol persisted; dataset properties without evidence marked needs_sequence_audit
SOURCE	agents/deepseek_meeting/metric_agent.md · deep_research_literature_agent.py:3354-3369

P03 Critic audit and Chief memory update

SYSTEM	Audit scope, provenance, reproducibility, leakage and unsupported certainty. The Chief must reconcile proposals, criticisms, rebuttals and dissent while preserving the cumulative candidate pool and explicit reasons for any downgrade.
USER	Adjudicate every dataset and model coverage candidate, check continuity with prior memory, then write strict JSON for the long-term state. Historical memory is an evidence anchor rather than a fixed recommendation template.
VERIFIED OUTPUT	literature_deep_research_memory.json/.md and screening decision CSV refreshed
SOURCE	agents/deepseek_meeting/critic_agent.md · chief_agent.md · deep_research_literature_agent.py:3374-3474