

# Blinded 50-round metric weighting and model ranking

## P12 Independent specialist weight proposals

|                 |                                                                                                                                                                                                                                          |
|-----------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| SYSTEM          | You are one role in a blinded multi-Agent meeting selecting metric weights for AMP binary classification. Never optimize weights for a named model, a desired Top-3 or a leaderboard position.                                           |
| USER            | Literature, Statistics and Screening Agents independently propose all eligible metric weights from literature consensus, anonymous dataset profiles and metric-level evidence. Distinguish literature, benchmark evidence and LLM prior. |
| VERIFIED OUTPUT | Three role-specific proposal streams retained for the initial meeting and every round                                                                                                                                                    |
| SOURCE          | agents/weight_meeting/shared_system.md · literature_agent.md · statistics_agent.md · screening_agent.md                                                                                                                                  |

## P13 Independent methodology audit

|                 |                                                                                                                                                                                                          |
|-----------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| SYSTEM          | Audit unsupported evidence, hidden model preference, test-set tuning, excessive metric dominance, redundancy, ignored calibration, prevalence and unresolved disagreement. Do not produce final weights. |
| USER            | Review all three blinded proposals and return criticisms, required changes, preferred metric directions, a leakage check and unresolved risks in strict JSON.                                            |
| VERIFIED OUTPUT | reviewer_agent_audit.json plus per-round Reviewer records preserved                                                                                                                                      |
| SOURCE          | agents/weight_meeting/reviewer_agent.md · llm_agent_weight_meeting.py:457-470                                                                                                                            |

## P14 Chief bounded consensus

|                 |                                                                                                                                                                                                                   |
|-----------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| SYSTEM          | Reconcile the expert proposals and Reviewer audit into the only accepted weight vector. Respond to criticisms and preserve disagreement and uncertainty.                                                          |
| USER            | Return all exact metric keys with weights in [0.005, 0.35], sum=1, no model-specific priority and at most 0.30 L1 change from the previous accepted vector. Runtime code validates and repairs schema violations. |
| VERIFIED OUTPUT | chief_initial_decision.json and round_001.json through round_050.json                                                                                                                                             |
| SOURCE          | agents/weight_meeting/chief_agent.md · llm_agent_weight_meeting.py:473-497                                                                                                                                        |

## P15 Fifty-round evidence aggregation

|                 |                                                                                                                                                                                                    |
|-----------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| SYSTEM          | The accepted vector scores every eligible model against the same normalized metric evidence. Model identities remain absent from weight-setting prompts, and the complete audit trail is retained. |
| USER            | Repeat bounded evidence updates for 50 rounds, calculate each model's weighted score, then aggregate median score, IQR, median rank and Top-3 frequency without a model-specific bonus.            |
| VERIFIED OUTPUT | 50 rounds · 600 weight rows · 900 model-round scores · Top-3: pepnet_standard / amplify_imb / C_AMPs-predict                                                                                       |
| SOURCE          | codex_agent_weight_meeting.py · iterative_weight_meeting.py · data/results_manual/codex_agent_weight_meeting/                                                                                      |