

## Evidence search, recommendation and persistent memory

Role-specific Agents retrieve, compress, challenge and reconcile literature evidence before freezing model, dataset and metric recommendation tables.

|        |                  |                |                    |
|--------|------------------|----------------|--------------------|
| 2,365  | 304              | 241            | 20                 |
| papers | evidence batches | compact chunks | recommended models |


### Agent contracts and responsibilities

versioned Markdown or control files

| ROLE     | DEFINITION FILE              | RESPONSIBILITY                                   |
|----------|------------------------------|--------------------------------------------------|
| Compress | evidence_compressor_agent.md | Retain traceable evidence while reducing context |
| Extract  | info_extractor_agent.md      | Produce structured paper and repository metadata |
| Scout    | model_dataset_agent.md       | Propose models and benchmark datasets            |
| Metrics  | metric_agent.md              | Define endpoints and evaluation evidence         |
| Critic   | critic_agent.md              | Challenge unsupported or incomplete claims       |
| Chief    | chief_agent.md               | Reconcile decisions and write persistent memory  |

### Execution sequence

ordered file handoffs



### Persisted evidence and stage handoff

project-state artifacts

| ARTIFACT                                   | ROLE                | AUDIT / HANDOFF MEANING                           |
|--------------------------------------------|---------------------|---------------------------------------------------|
| evidence_pool.json                         | raw evidence        | 2,365 papers / 304 batches                        |
| compact_evidence_pool.json                 | compressed evidence | 2,361 papers / 241 chunks                         |
| literature_meeting_screening_decisions.csv | audit decisions     | 487 identities / 59 accepted                      |
| literature_deep_research_memory.json/.md   | long-term memory    | machine-readable state plus human-readable record |
| recommended_{models,datasets,metrics}.csv  | frozen handoff      | 20 models / 3 datasets / 44 metric rows           |

#### Memory rule

The next literature run reads and updates the stored memory; the recommendation list is versioned evidence, not a permanently fixed template.