

Fifty bounded rounds convert metric evidence into consensus ranking

Three specialist proposals are reviewed independently and reconciled by a Chief Agent before each accepted metric vector scores all models across datasets.

12	50	600	900
eligible metrics	decision rounds	weight rows	model scores

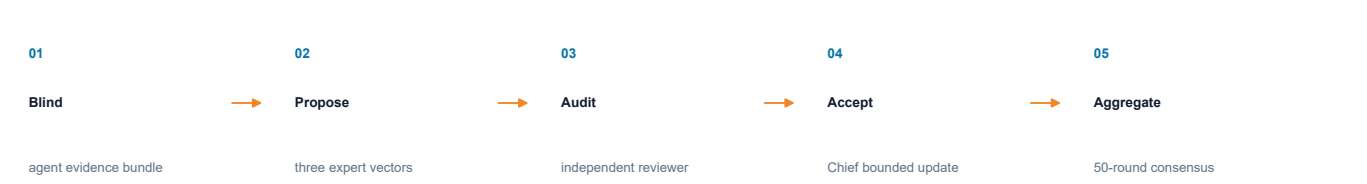
Agent contracts and responsibilities

versioned Markdown or control files

ROLE	DEFINITION FILE	RESPONSIBILITY
Literature	weight_meeting/literature_agent.md	Judge biological relevance and evidence support
Statistics	weight_meeting/statistics_agent.md	Judge validity, uncertainty and metric behavior
Screening	weight_meeting/screening_agent.md	Judge operational robustness and screening utility
Reviewer	weight_meeting/reviewer_agent.md	Audit proposals against blinded evidence
Chief	weight_meeting/chief_agent.md	Accept a bounded weight vector or request revision
Advisor	weight_meeting/research_advisor.md	Interpret stability, complementarity and Top-3

Execution sequence

ordered file handoffs



Persisted evidence and stage handoff

project-state artifacts

ARTIFACT	ROLE	AUDIT / HANDOFF MEANING
agent_evidence_bundle.json	blinded evidence	3 datasets / 18 models / eligible metrics
*_agent_proposals.json	specialist proposals	Literature, Statistics and Screening vectors
reviewer_agent_audit.json	independent audit	proposal-level challenges and decisions
chief_initial_decision.json + rounds/round_001-050.json	decision state	initial vector plus every accepted update
weights/scores/ranking_50_rounds.csv	terminal evidence	current Top-3: pepnet_standard / amplify_imb / C_AMPs-predict

Aggregation equations

Single round: $S_m^{(t)} = \sum_k w_k^{(t)} q_{mk}$; terminal evidence: $\hat{S}_m = 50^{-1} \sum_{t=1}^{50} S_m^{(t)}$. No model-specific bonus is applied.